

**Weekly Oxford Worldwide**

DEPARTMENT FOR  
CONTINUING  
EDUCATION




# **PYTHON PROGRAMMING FOR DATA SCIENCE – PART 2**

## **MASSIMILIANO IZZO & NICHOLAS DAY**

### **LECTURE 4**

### **SUPERVISED LEARNING: CLASSIFICATION WITH SCIKIT-LEARN**



# This week

- Classifiers
- Performance metrics for Classification
- Classification using MNIST dataset

# Classification

- Classification tasks can be:
  - binary (two classes, generally coded a 0 and 1)
  - multi-class
  - multi-label
- Performance metrics are trickier for classification. Accuracy would be the most intuitive obvious one to measure

$$accuracy = \frac{\# \text{ correctly predicted records}}{\# \text{ total records}}$$


# Logistic Regression

Binary classification (classes 0 and 1)

Estimates the probability that a sample belongs to a certain class by training a (linear) regressor that will return scores in the  $(-\infty, +\infty)$  interval

then passing the output of the regressor to a **logistic (sigmoid) function**

The output will be a value between 0 and 1.  
If the output is  $> 0.5$  assign to class 1  
otherwise assign to class 0



$$\hat{p}(\boldsymbol{\theta}) = \hat{p}(\mathbf{w}, b) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-(\mathbf{x}^T \mathbf{w} + b)}}$$

$z$  and  $\hat{p}$  are the regressor score and the predicted probability for the positive class ( $y = 1$ )

# Logistic Regression: the log loss cost function

We can train a Logistic Regression classifier using Gradient Descent, but we cannot use the MSE as a cost function to minimise


$$c(\boldsymbol{\theta}) = \begin{cases} -\log \hat{p} & \text{if } y = 1 \\ -\log(1 - \hat{p}) & \text{if } y = 0 \end{cases}$$

where  $y \in \{0, 1\}$  is the actual label value, and  $\hat{p}$  is the predicted probability of positive instances ( $y = 1$ )

The overall cost function over a training set with  $m$  samples is:

$$\text{logloss} = J(\theta) = -\frac{1}{m} \sum_{i=1}^m y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)$$

# Log-loss cost function for class 1





# Softmax Regression

Extension of the logistic regression to multi-class scenarios.

The regressor will return a regression score  $z_k$  for each class  $k$

Rather than using the sigmoid we use the softmax function:

$$\bullet \hat{p}_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

$$\bullet \sum_i^K \hat{p}_i = 1$$

$K$  = number of classes

$k$  = current class

$z_k$  = regressor score for class  $k$

$\hat{p}_k$  = probability for class  $k$

The class with highest probability will be the predicted class of the softmax regressor

# Softmax Regression: the cross-entropy cost function

The cost function used for softmax regression is a multi-class extension of the log-loss function, called cross-entropy

$$\text{crossentropy} = J(\boldsymbol{\theta}) = -\frac{1}{m} \sum_{i=1}^m \sum_k^K y_k \log \hat{p}_k$$

$m$  = number of records/samples in training set

$K$  = number of classes

$k$  = current class


$z_k$  = regressor score for class  $k$

$\hat{p}_k$  = probability for class  $k$  (this is a function of  $\boldsymbol{\theta}$ )



# Support Vector Machines

A support vector machine (SVM) constructs a hyper-plane or set of hyper-planes in a high or infinite dimensional space, which can be used for classification, regression or other tasks. Intuitively, a good separation is achieved by the hyper-plane that has the largest distance to the nearest training data points of any class (so-called functional margin), since in general the larger the margin the lower the generalization error of the classifier

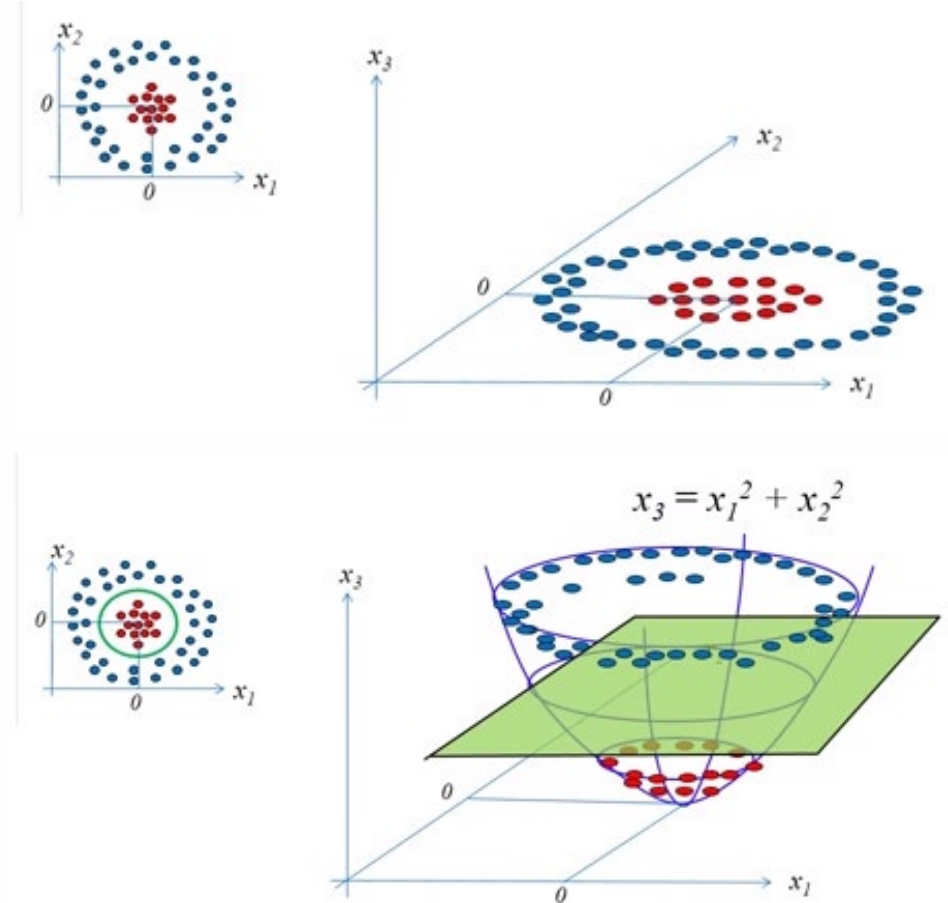


# SVMs: the Kernel Trick

To solve a nonlinear problem with SVM:

1. We transform the training data onto a **higher dimensional feature space** via a mapping function  $\phi$ .
2. We train a linear SVM model to classify the data in this new feature space.
3. Then, we can use the same mapping function  $\phi$  to transform unseen data to classify it using the linear SVM model.

The **kernel trick** avoids the explicit mapping that is needed to get linear learning algorithms to learn a nonlinear function or decision boundary.




# Online SVMs with the hinge loss

$$J(\theta) = J(\mathbf{w}, b) = \underbrace{\frac{1}{2} \mathbf{w}^T \mathbf{w}}_{\text{regularization term}} + C \underbrace{\sum_{i=1}^m \max(0, 1 - t_i(\mathbf{w}^T \mathbf{x}_i + b))}_{\text{hinge loss term}}$$

Traditional SVM are trained offline (batch-training)

However, we can train online SVMs using gradient descent, just like we train logistic or softmax regression classifiers


Rather than using the log loss we use the **hinge function** as our cost function



## Stochastic Gradient Descent Classification

<https://scikit-learn.org/stable/modules/sgd.html>

- The class `SGDClassifier` implements a plain stochastic gradient descent learning routine which supports different loss functions and penalties for classification.
- The default function scikit-learn is the Linear SVM decision function



[https://scikit-learn.org/stable/auto\\_examples/linear\\_model/plot\\_sgd\\_separating\\_hyperplane.html](https://scikit-learn.org/stable/auto_examples/linear_model/plot_sgd_separating_hyperplane.html)

# Nearest Neighbours Classification

- a type of *instance-based learning* or *non-generalizing learning*
- Classification is computed from a simple majority vote of the nearest neighbours of each point
- KNeighborsClassifier
- RadiusNeighborsClassifier

# Naive Bayes

- Naive Bayes methods are a set of supervised learning algorithms based on applying Bayes' theorem with the “naive” assumption of conditional independence between every pair of feature records given the value of the class variable.

$$\hat{y} = \underset{y}{arg \max} P(y) \prod_{i=1}^N P(x_i|y)$$

- The different naive Bayes classifiers differ mainly by the assumptions they make regarding the distribution of  $P(x_i|y)$




# Decision Trees

- Decision Trees (DTs) are a non-parametric supervised learning method used for classification (and regression). The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features.
- Advantages:
  - Simple to understand and to interpret. Trees can be visualised.
  - Requires little data preparation (no normalization)
  - Able to handle both numerical and categorical data.
- Disadvantages:
  - Overfitting
  - Learning the optimal DT is an NP-complete problem => heuristics

# Neural Networks: Multi-layer Perceptron

- To be seen...



# Ensemble methods

- The goal of **ensemble methods** is to combine the predictions of several base estimators built with a given learning algorithm in order to improve generalizability / robustness over a single estimator
  - Averaging methods: Random Forests
  - Boosting methods: Ada Boost, Gradient Tree Boosting
  - Voting Classifier


# Performance Metrics: Confusion Matrix

$$accuracy = \frac{\text{correct predictions}}{\text{total predictions}} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

$$F1 = 2 \frac{precision \times recall}{precision + recall}$$



How many selected items are relevant?



How many relevant items are selected?



# Performance Metrics: Area under the ROC Curve

$$TPR = recall = \frac{TP}{TP + FN}$$

$$FPR = \frac{FP}{FP + TN}$$

